

1 Intellectual Merit

1.1 Aims

As phylogenetic and functional approaches transcend species-based estimates of diversity, we seek new answers to the age-old question [1], “What limits diversity?” At the intersection of evolutionary history, environmental opportunity, functional capacity, and genetic adaptive potential, a suite of fundamental processes (drift, migration, mutation, speciation, competition, environmental filtering) and emergent properties (e.g. community structure, ecosystem productivity) lead to largely unexplored feedbacks between species, functional, phylogenetic, and population genetic diversity. Thanks to initiatives such as NSF’s Dimensions of Biodiversity (DoB) Program, we have gained enough insight into feedbacks among these multiple dimensions of biodiversity to know that they complicate previously simple interpretations of single dimensions. For instance, traits such as body size may directly shape responses to climate change by impacting population demography, thus driving biogeographic patterns of lineage distribution, which themselves define the conditions to which lineages adapt [2, 3]. However, we have yet to add complexity to our predictive models of biodiversity and thus not gained a mechanistic understanding of the interactions among diversity dimensions. There is now an unprecedented opportunity to advance understanding of basic and emergent interactions among dimensions, as rich data describing detailed and large-scale multi-dimensional data become available through DoB and other similar initiatives.

Because measurements of diversity reflected in the scale of DoB projects range from small archipelagos to whole biomes, we propose to synthesize and integrate across projects to understand the general rules governing biodiversity. Making sense of diverse and multi-dimensional data within and between such diverse biogeographical systems requires (1) having the **informatic ability** to synthesize and share them and (2) the **quantitative tools** to extract insight from synthesized data about the complex, multi-dimensional processes underlying biodiversity patterns. **We propose to fill these two gaps by (1) building an open source platform to manage and query diverse streams of biodiversity-relevant information, and (2) developing a stochastic model of biodiversity that incorporates a range of mechanistic processes that jointly predicts species abundance, trait distributions, phylogenetic structure, and population genetic diversity.** Through this model we will be able to quantify the contribution of three distinct classes of processes to the structure of biodiversity across disparate systems: (i) Stochastic immigration, speciation and extinction, (ii) species interactions with the biotic and abiotic environment, and (iii) historical contingencies.

To understand the universality of these processes, we will (3) test our model across five biogeographic systems for which multi-dimensional biodiversity data are now available, due largely to DoB efforts. They include two naturally insular systems with differing degrees of isolation, age, and adaptive radiation (the Hawaiian island chronosequence and the marine lakes of Palau) and three forest systems with key gradients of diversity and evolutionary history, spanning scales of historic continental disjunctures (the Nearctic forests of North America and Asia), entire biomes (Amazonia), and large regions of unique endemism (the coastal Atlantic Forest of Brazil; Fig. 1). The two island projects additionally afford an explicitly temporal test of our model’s dynamics—via the sediment cores from Palau’s marine lakes and the chronosequence in Hawaii. By selectively augmenting data gathered by these DoB projects with new field observations, we will be able to test our model and begin to verify the universality of processes regulating the diversity of life. The new data we collect will add substantially to the Tree of Life, particularly for under-sampled clades such as arthropods and marine invertebrates.

1.2 Background

Understanding the generation and maintenance of biodiversity is critical to manage and mitigate the effects of anthropogenic change on the diversity of living systems; without intervention, recovery from the current crisis [4] could take tens of millions of years [5]. Despite the lack of concrete answers to what limits diversity, decades of research have generated important insights into the processes driving diversity across multiple dimensions in specific systems. Still, mechanistic hypotheses from these studies have yet to be quantitatively and jointly tested across systems with a rigorous framework where their universality can be evaluated.

Current eco-evolutionary theory is not up to this task. Quantitative models inspired by the Equilibrium Theory of Island Biogeography [6] and the Neutral Theory of Biodiversity and Biogeography [7], for instance, place emphasis on the equilibrium between rates of immigration, speciation, and extinction. Yet while empirical work on islands and latitudinal gradients have demonstrated the significant role of immigration, speciation and extinction histories [8, 9], these neutral theories of eco-evolutionary drift are limited. For example, interspecific interactions and the environment have been shown empirically to be important in determining patterns such as the latitudinal diversity gradient [10–12]. Moreover, neutral drift theories do not accurately capture temporal dynamics [13, 14] and ecological interactions have been flagged as potential drivers of the fast dynamics observed in real communities that neutral theories fail to accurately predict [13, 15]. Conversely, models based on detailed ecological interactions may be so over-parameterized as to render them uninterpretable [16]. Current theory fails to predict multiple dimensions of biodiversity, which can lead to very different processes producing indistinguishable predictions for single dimensions [17–20].

The perspectives from both neutral-like eco-evolutionary drift and ecological interaction-based coexistence and diversification are typically applied in equilibrium [7, 21–23]. However, the history underlying a system at any given point might be such that the system is far from an equilibrium state –something increasingly relevant in the Anthropocene [24]. Examples of this non-equilibrium hysteresis include island biogeography [25], differences in otherwise similar biomes across East Asia and North America [26, 27], and communities assembled from refugia after climate change [28–30].

Evaluating how general these observations are across systems, and how they influence species-level, phylogenetic, population genetic, and functional trait diversity, necessitates the collection, management and analysis of multi-dimensional data on i) the abundance of species, ii) their geo-referenced occurrences, iii) ecologically-relevant traits, iv) phylogenetic and population genetic histories, and v) the environment. We propose to build such informatic synthesis tools and to quantitatively compare the support for these different hypothesized mechanisms across five systems.

1.3 Informatic platform for dimensions of biodiversity data

Large scale, system-based collaborative projects - including many funded by the DoB program - often have difficulty making the best use of the volume and types of data they generate. Data management strategies are usually developed independently by each team or laboratory, often varying across teams. Members of collaborating laboratories frequently work on their own locally stored datasets, with local domain expertise and data administration and curation. Smaller or more narrowly focused projects can afford to absorb the administrative and cognitive overhead of fully

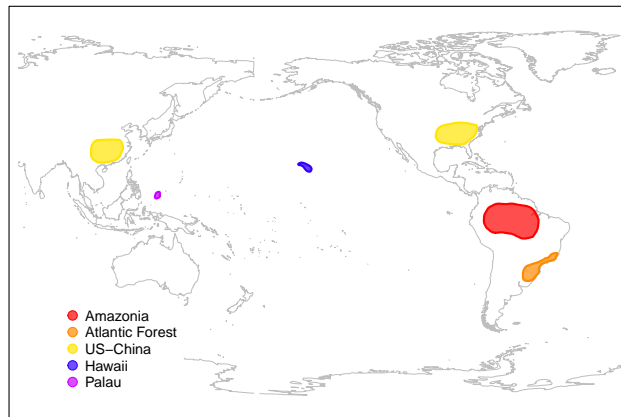


Figure 1: Map of the included DoB projects showing their breadth of regions and scales.

integrating data platforms and management. However, larger projects - both in scope and scale - are significantly negatively impacted by this. The problem is especially acute for projects with multiple PIs distributed across the globe (e.g. US-China and US-Brazil partnerships), as well as those generating vast quantities of heterogeneous data.

Various isolated repositories do exist for warehousing some of these data independently. Initiatives such as iDigBio, GBIF, Dryad, GloBI, and NCBI provide some facility for depositing, cataloguing and accessing occurrence records, sequence, morphological and ecological data, and have been extensively used by the many DoB projects funded to date. However, the linkages that would crucially support re-integration of these data are in nearly all cases missing, which limits the ability to derive broader understanding across projects. While all the mentioned repositories are essential services to the global scientific community, they each have internal missions that are divergent from the needs of biodiversity project teams for data integration. Further, there are still limited metadata standards that are both general and comprehensive for ecological datasets, hampering integration, although new approaches have been proposed [31]. We argue that simply depositing data into such repositories is not enough. A key step forward is building a framework where data resources are standardized enough to simplify the process of feeding the right types of data into modeling frameworks. To better understand biodiversity pattern and process, DoB projects and other biodiversity research teams should be able to ask questions about the multivariate forces that influence multiple systems; thus there is a critical need to assure that datasets talk to each other. **We will build a proof-of-concept platform that supports storing data for broadest possible use and also enhances modeling frameworks, such as those discussed below.** The outcome will be a scalable model knowledgebase that crosscuts and connects all dimensions of spatiotemporal biodiversity data.

1.4 Theory to integrate dimensions of biodiversity

Recent efforts by Hickerson and his team have jointly modeled species abundance and community-level genetic variation through time [32]. Their work explicitly links a forward-time neutral model of community assembly with a backward-time neutral coalescent model of population genetics, fitting the model to data with approximate Bayesian computation [ABC; 33]. **We propose to extend this model to jointly predict species abundance, phylogeny, population genetic, and functional trait distributions within ecological communities (Fig. 2).** To achieve this theoretical advance, we will incorporate speciation [34] and trait evolution to existing work in neutral theory [22, 35]. By including population genetic and phylogenetic information when estimating the model and assessing its fit to data, we will capture signatures of temporal dynamics and scrutinize our work exactly where neutral theory fails most noticeably [13]. In addition to modeling genetic axes of biodiversity data under neutrality, we will extend the neutral model to incorporate non-neutral processes, such as density dependence [36], and the ability of fitness to evolve within and between species [15], all of which will help to reconcile modeled temporal dynamics with reality. We will link these coarse-grained processes, for which mathematical details have been well-studied, to our trait data. For that, we will use model parameters that capture trait-based interactions with other species [37, 38] and environments [39–41]. Table 1 shows how these processes enter into our model.

We will fit this joint model of multiple biodiversity dimensions (species abundance, phylo- and population genetic, and trait) to empirical data collected across the five DoB projects using both ABC and new machine learning approaches [42]. Evaluating how well the model fits empirical data across systems will allow us to understand whether a parsimonious set of processes encapsulates pervasive patterns in the diversity of life. Quantifying how the specific best-fit parameter values of the model change across systems will allow us to understand the contribution of these different

Process	(symbol)	Neutral	Environment x Trait	Trait x Trait
birth	(b)	equals fitness	equals fitness	positive function of fitness and negative function of competition
death	(d)	constant	constant	constant
immigration	(ν)	constant	positive function of potential fitness in colonizing environment	negative function of potential competition in colonizing community
fitness	(r)	constant	positive function, depending on coefficient ρ , of environmental suitability of trait	constant
environmental filtering	(ρ)	0	strength of environmental filtering	0
competition	(a_{ij})	0	0	positive function of trait-overlap
speciation	(λ)	constant	positive function of environmental suitability of incipient trait	negative function of trait-overlap of incipient species
mutation	(μ)	constant	constant	constant
trait evolution	(σ, θ)	Brownian	Ornstein-Uhlenbeck	Ornstein-Uhlenbeck or Brownian

Table 1: Model parameters.

processes to the structuring of biodiversity across disparate systems, and spatial and temporal scales.

2 Proposed Research

2.1 Research objectives and hypotheses

We have three objectives:

R1 To design and implement an informatics platform that can synthesize and facilitate the sharing of diverse data streams associated with DoB projects and similar biodiversity initiatives. We will make this platform open source and usable both through a graphical user interface and an R [43] package. We will implement it across the five DoB projects represented by co-PIs on this current proposal (Hawaii: Rominger and Gillespie; Palau: Dawson; US-China: Soltis; Amazonia: Guralnick and Owens; Atlantic Forest: Carnaval and Hickerson).

R2 To build a joint model for the multiple dimensions of biodiversity—species abundance, genetic, and functional. This model will build off on-going work (by Rominger, Hickerson, Harmon, and Chase) to combine macroecological predictions with phylo- and population genetic measures. Model fitting methods based on ABC and machine learning will be created. The entire modeling framework will be released as open source software.

R3 To test this joint model across five diverse systems for which data has been synthesized by the informatics platform from existing DoB sampling efforts and selective resampling. We have included in this proposal five mature DoB projects whose data represent compelling unit tests for our informatics platform, and whose underlying evolution and ecology will test the generality of our model. We have identified data types (e.g. abundance, genetic and/or trait data) where new targeted sampling within these projects will greatly improve our ability to test our model and document the

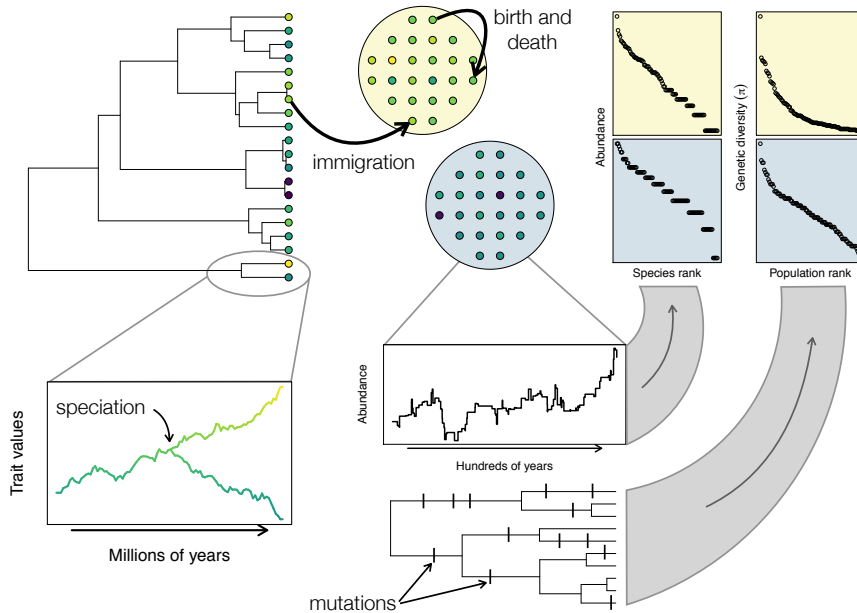


Figure 2: The processes captured by our model and the data types predicted. Colors correspond to trait values. The processes are: birth, death, immigration, and mutation in a coalescent framework (all of which operate primarily on short time scales), and speciation and trait evolution (all of which operate primarily on longer time scales). These processes may be modulated by trait-based interactions among species and between species and their environments as presented in Table 3. Different environments and their trait compatibility are illustrated for two hypothetical communities by the lighter background colors. The modeled processes generate the four data types predicted by the model: species abundance, population genetic diversity across species, phylogenetic relationships between species, and trait distributions across species.

raw diversity of life.

The modeling approach we develop in **R2** allows us to address five hypotheses central to ecology and evolution at an unprecedented quantitative caliber and grounded in data aggregated (**R1**) from incredibly diverse systems (**R3**). Our approach casts these complex hypotheses as clear correlations (see Fig. 3) between the parameters of our model (Table 1) when fit to data across systems. Because the model encapsulates a range of processes and jointly predicts many data types, interrogating its behavior represents a rigorous, multi-faceted test of hypotheses [18–20].

H1: Isolation fosters diversification. This has long been assumed, documented empirically [44, 45], and demonstrated theoretically [35]. However, it remains to be assessed whether isolation is key to the striking examples of endemism and radiation at the species, genetic, and functional levels, or if it is only secondary to other processes such as strong ecological interactions. Looking across systems, this hypothesis predicts a negative correlation between estimated immigration (ν) and speciation (λ ; Fig 3).

H2: Environmentally-driven speciation is important across systems, but depends on intrinsic dispersal ability of taxa. Whether environmental gradients can promote diversification remains an open debate [11, 46, 47]. Our model can test whether including interactions between traits and environments is necessary to predict patterns of species abundance, phylo- and population genetic diversity, and trait distributions. After accounting for the effect of H1, this hypothesis predicts a positive correlation between estimated speciation (λ) and the measured environmental heterogeneity of a study region (Fig. 3), which we can quantify spatially and temporally for many

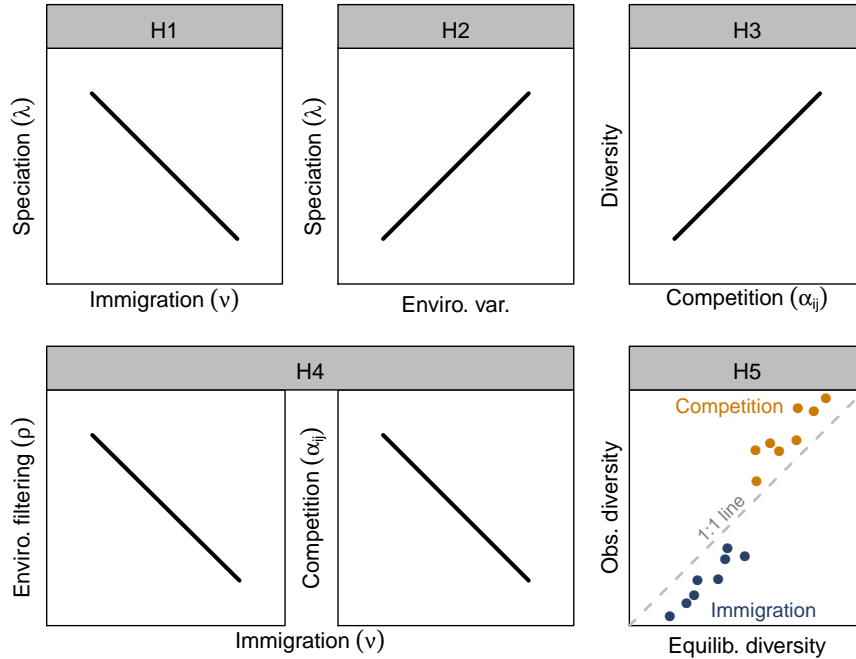


Figure 3: Correlations between inferred parameters (see Table 1) of our model under the assumption that our five hypotheses are true. In **H5**, “competition” and “immigration” refer to communities inferred to be competition-driven versus immigration-driven, respectively.

studies.

H3: High diversity systems result from increased strength of species interactions. This is a classic hypothesis in the latitudinal diversity gradient literature [12] and questions of ecosystem complexity and stability [48, 49]. If we find that high diversity systems are systematically better-fit by models with strong trait-based species interactions, this lends strong support to the hypothesis. Across all systems, this hypothesis predicts a positive correlation between diversity measured across all biodiversity dimensions and the magnitude of estimated trait-based competition (α_{ij} ; Fig. 3).

H4: Assembly by immigration (versus *in situ* speciation) leads to weaker species interactions and thus more apparently neutral distributions of species abundances, genetic and functional diversity. If high diversity systems are the result of strong species interactions, do these interactions arise more readily in communities assembled by immigration or *in situ* speciation? We posit the latter, due to increased opportunity for co-evolutionary arms race-like dynamics [50, 51]. Conversely, immigration-assembled communities will draw from pools of the best, often most generalist, colonists and thus appear effectively neutral. Thus, this hypothesis predicts a negative correlation between inferred immigration and the magnitude of parameters capturing purely non-neutral processes (i.e. trait-mediated interactions, both species-species [α_{ij}] and species-environment [r]; Fig. 3).

H5: Most systems have not reached an equilibrium in diversity. Thus historical contingencies are critical for understanding contemporary diversity. We posit that the mechanistic importance of the previous four hypotheses is heightened by systems being out of equilibrium. For systems in equilibrium, dynamical processes can no longer be inferred, or their effect size will be diminished [52]. The time to equilibrium, and whether or not our model has reached it, is something we can quantify. We can then continue running the model until equilibrium is reached. Across systems we predict discrepancies between observed diversity across all dimensions and the modeled equilibrium diversity (Fig. 3). Whether this discrepancy is toward higher or lower observed diversity than equilibrium yields interesting additional and complementary hypotheses that we will

explore:

H5a Systems constrained by immigration [e.g. arthropod communities on young substrates in Hawaii; 53] will show observed diversities lower than equilibrium.

H5b Systems driven by strong interactions will show diversity overshoots in which non-equilibrium diversity is higher than equilibrium diversity because interactions initially drive diversification beyond the intrinsic carrying capacity of a system [54].

2.2 Methods

2.3 Informatics pipeline

Existing informatic platforms and metadata standards now only manage subsets of the data we envision synthesizing: abundance and/or occurrence, trait, genetic, and environmental. We identified major barriers to creating this multi-dimensional biodiversity platform for projects that span biogeographic regions as part of an NSF-funded workshop on DoB data management (PIs Carnaval and Soltis). These barriers are (1) allowing for short-term data entry and long-term data storage and sharing; (2) integrating abundance, phylogeny, and environmental data into a data management environment primarily intended for specimen data; and (3) educating researchers about proper data management using principles from data science. This last barrier we address with substantial curriculum development, detailed in section 3.1 on training and Broader Impacts.

It is critical to insure that trait, occurrence, and abundance data are standardized properly across all projects. We will use entity-quality relationships to standardize key traits shared across the tree of life, namely body size and metabolic rate. Those traits specific to groups will also be modeled using a similar approach. Both individual and species-level traits will be synthesized following a new relational data model developed by Co-PI Guralnick. Abundance data will be standardized using Humboldt Core [31] which captures necessary inventory metadata. Efforts at standardizing content will assure both model-ready data, and reproducibility and re-use for others.

We will build an open source R package that allows users to manage and synthesize their data via three cloud-based and/or open repositories: (1) Google Sheets for data entry and short-term curation; (2) Amazon Relational Database Services (RDS) for mid-term curation; and (3) open repositories detailed below for long-term curation.

Data for such large initiatives as DoB projects necessitates that many participants be involved in the collection and curation of data. Cloud-based platforms are ideal for distributed data collection because they can be accessed remotely and track the version history produced by multiple users. To this end, Rominger, working with undergraduate students in focused mentoring associations, has developed an open source R package [55] that facilitates the entry and error checking of data using R and Google Sheets, either through custom scripts or a graphical user interface (GUI) built with Shiny [56]. This approach will be extended to facilitate the entry of data types including abundances, phylogenies, and raw sequence data. For data that cannot be easily recorded in tabular format (e.g. phylogenies), links to those files will be recorded instead.

Once data entry is complete we will house data on Amazon RDS and provide functionality in our R package and Shiny app for data transfer to Amazon RDS. Amazon RDS provides a scalable solution that reduces administration costs and provides a simple solution for access to content via query APIs. It is also easy to connect to such databases via R [57]. We will assure open access to resources via R and the Shiny app to team members and the larger community, thus hastening increased re-use.

Solutions for open, long-term curation of all data types of interest to our project have already been pioneered. The missing component is practical software to link data from these repositories for a given project. We will fill this gap by making our extended R package work not only with Google Sheets, local files and Amazon RDS, but with the APIs of data aggregators and curators. We will

specifically target GBIF for occurrence data [58], NCBI for sequence data [59, 60], OpenTree for phylogenies [61], Map of Life for abundance data [31, 62], TraitBase for traits [63], and GloBI for biotic interactions data [64]. Importantly, all these curation platforms already have dedicated R packages [65–69] that interface with their APIs, which we will leverage for our newly developed package. We will use these existing resources to create high-level functions that facilitate the publishing of data from Google Sheets, local files, and Amazon RDS to these open repositories. Again, this functionality will be made available both as a standalone R package and as a Shiny GUI.

To facilitate our specific data needs for testing our model of eco-evolutionary community assembly, we need to be able to compile data by three potentially overlapping ways, each targeting specific shared attributes across disparate data types to merge—or “join”—relevant information from all data types. We will use a PostgreSQL [70] backend in our R package to allow for these three joins:

1. Spatial joins: we must be able to collect all relevant data, from occurrences, abundance, and environments, for a specific location or region.
2. Phylogenetic joins: we must be able to collect all relevant data for a clade defined by a sample of its constituent members
3. Data availability joins: we must be able to identify clades and regions for which sufficient data are available and then collate those data.

These joins will allow us to, for example, find all taxa that have abundance data at a given location, query their trait values which may have been measured from specimens at different locations, compile a phylogenetic tree for those taxa and the files containing their raw sequence data.

2.4 Collection of new data

Because we selected five mature DoB projects to populate our bioinformatics pipeline and test our model, many data resources are already available; our new data collection efforts will focus on filling specific gaps to ensure continuity of data types across projects. Table 2 summarizes available data and gaps to be filled.

Project	Taxon	Abund.	Occur.	Pop. gen.	Phylo.	Traits	Enviro.
Hawaii	Arthropods						
Palau	Algae						
Palau	Invertebrates						
Palau	Fish						
Atlantic Forest	Plants						
Atlantic Forest	Butterflies						
Atlantic Forest	Herpetofauna						
Amazonia	Plants						
Amazonia	Birds						
US-China	Plants						

Table 2: Data resources available, and planned collection of new data. **Blue cells** correspond to sufficient existing data, **tan cells** correspond to data gaps that will not be filled, **black cells** represent new data to be collected by this proposed project.

Gap-filling field work will complement our abundance, phylogenetic and population genetic, and trait sampling. Our primary source for choosing sites will be the experience of the PIs and the knowledge of local collaborators, supplemented with information from satellite images and remote sensing data. Below we detail the status of specific projects, and what new data are needed to build robust phylogenies, document the multiple dimensions of biodiversity, and test our model.

Hawaii. This project uses the dynamic geomorphology of the Hawaiian Islands to examine the evolution of entire communities of arthropods over extended time. A combination of community ecological approaches with genomics of select evolutionary lineages allows us to determine the importance of changing functional roles of taxa within communities as they differentiate and assemble, and the role of that dynamic community in fostering diversification. Currently missing from this effort, and necessary for the model, are complete data on functional traits and phylogenetic relationships. Body size data already exist for $\approx 10^6$ specimens and trophic interactions exist for a much more limited subset, namely herbivorous hemipterans [53]. We will add trophic interactions for all herbivorous insects following the methods in Rominger et al [53], as well as more general trophic position (e.g. predator, parasitoid, etc.) for all arthropods. In addition, collaborator Boettiger is compiling trophic interaction data for the Hawaii DoB based on field observations and sequencing of gut contents. Genetic resources are quickly accumulating for the Hawaii project following novel new metabarcoding practices [71]. We will augment these data, which provide sequences from a single marker for all specimens, with multi-locus sequencing for targeted voucher specimens in order to construct a robust tree for ≈ 2000 Hawaiian arthropod species (≈ 100 of which we will add through this grant). This will be the first time such a mega tree has been constructed for Hawaiian arthropods, many of which have never been phylogenetically treated at all.

Palau. Extensive abundance and trait data (including body size, trophic position and habitat preference) have already been collected. This system also has an aggressive COI barcoding initiative covering all invertebrates found (459 vouchered species and subspecies in 205 genera). We will again use multi-locus sequencing of targeted vouchers representing ≈ 100 genera to build a comprehensive phylogeny for invertebrates in Palau's unique marine lakes, augmenting ≈ 1200 barcoded samples. We will additionally analyze ancient COI barcode eDNA using myBaits target capture from down-core sediments followed by Illumina HiSeq sequencing to reconstruct community composition and dynamics including the immigration or origination of new alleles through time. A total of 75 m of cores span the Holocene in 6 lakes, and are in storage at the LacCore facility (U. Minnesota).

Atlantic Forest. We will focus on amphibians (> 100 Hylidae tree frogs within Boana, and the Phyllomedusinae), lizards (within the Eublepharinae), butterflies (≈ 50 species in the Biblidinae subfamily and the satyroide clade), and plants (≈ 200 species within the Miconieae, 50 within the Bignoniaceae). For these groups, data on geographic occurrence, population genetics, phylogenetics, and functional traits already exist or can be straightforwardly gathered or enhanced with strong involvement of CUNY undergraduate students and the help of Brazilian collaborators. A backbone multi-locus phylogeny already exists for these groups, but we will add ≈ 100 each to vertebrates and plants to increase the completeness and resolution of those tree. To ensure completeness of the data table, we will generate barcode data for all available and yet unsequenced individuals of the target groups, following standard protocols [71–73]. For the amphibian data, a strong barcoding effort is in place already. The primary information gap in the Atlantic Forest are curated data quantifying abundance. To address this, we will (1) merge and curate abundance inventories already available for plants (a table of abundance of woody plants is available for > 20 sites), butterflies (7 sites, available through collaborator Freitas), amphibians and reptiles (5 sites, collaborator Rodrigues); (2) complete five replicate surveys to complement the vertebrate abundance data, covering both lowland and montane areas of the forest; and (3) compile data on the history of biocollections for these groups to build a model of occupancy from existing occurrence data. These spatially broad occupancy estimates can be combined with the spatially narrow but detailed abundance data and used in our modeling framework.

The primary information gap in the Atlantic Forest are curated data quantifying the abundance of amphibians (within a group of ≈ 100 Hylidae tree frogs within Boana, and the Phyllomedusinae), lizards (within the Eublepharinae), butterflies (≈ 50 species in the Biblidinae subfamily and the

satyroide clade), and plants (≈ 200 species within the Miconiae, 50 within the Bignoniaceae). For these groups, data on geographic occurrence, population genetics, phylogenetics, and functional traits already exist or can be straightforwardly gathered with strong involvement of CUNY undergraduate students, and the help of Brazilian collaborators. We will take a three-pronged approach to generate useful abundance data. The first is to merge and curate abundance inventories available for plants (a table of abundance of woody plants is available for 20+ sites, and can be obtained through collaborator Saiter), butterflies (7 sites, available through collaborator Freitas), and amphibians (2 sites, collaborator Rodrigues). The second is to complete five replicate surveys to complement the amphibian abundance data, covering both lowland and montane areas of the forest. We will also compile data on the history of biocollections for these groups to build a model of occupancy from existing occurrence data. These spatially broad occupancy estimates can be combined with the spatially narrow but detailed abundance data (to be collected or curated) and used in our modeling framework. To ensure completeness of the data table, we will generate barcode (mitochondrial or chloroplast) data for all available and yet unsequenced individuals of the target plant and butterfly groups (amphibian barcode data are available through collaborator Lyra). These sequence data will provide population genetic characterizations, and be grafted to multi-locus phylogenetic trees already available for these groups. Trait data are available for most species (Table 2), but will require literature search and curation.

Amazonia. We will focus on birds and two plant lineages—Lecythidaceae and Miconieae—for which the most complete data exist. For the plants, we will compile publicly-available abundance data from forest inventory plots [74–76] to augment the detailed occurrence, genetic, and functional trait data that have already been collected for this group. To make use of the many other species inventoried in forest plots, we will also target voucher specimens to generate new sequence data that can be added to existing resources for plant community phylogenetics [77, 78]. For birds, we will perform targeted mist-netting and point surveys to fill geographical and taxonomic gaps in existing abundance data from museum records, estimates from the eBird database, and long-term survey data provided by co-PI Robinson and collaborator Loiselle.

US-China. The US-China project focuses on levels of diversity across the classic eastern Asia-eastern North American floristic disjunction. Detailed sampling of plant genetic, trait, and phylogenetic diversity and soil and leaf microbial diversity at six sites in eastern North America and 5 sites in China will be available through the funded project; remote sensing data for 3 of the 6 North American sites are also currently available. Each site is linked to abundance data from NEON (US) and CERN (China) plots, with additional abundance data for woody plants provided by the Forest Inventory and Analysis Program (US Forest Service) and partner organizations in China. We seek to collect sequence data for plant species and soil microbes at four additional NEON sites in North America to improve sampling across the large geographic and climatic gradient and to add remote sensing data for additional NEON sites as the data become publicly available. We will add trait data from publicly available sources for any new species sampled. Occurrence data will be obtained from public databases. Phylogenies for 30 plant genera from across the disjunction will be completed in the funded project.

2.5 Theory development

Hickerson and his PhD student Overcast, who will be funded on the current proposal, have extended the island-mainland neutral model [79] to include coalescence of alleles under neutral mutation, thereby allowing joint predictions of species abundance distributions (SADs) and population genetic metrics such as genetic diversity (π). Rominger, Hickerson, Overcast, and collaborators Harmon and Chase will extend this work further to complete objective **R2** by including the process of speciation [34], the evolution of traits, and ecological interactions (between species, and between species and

their environments) based on those traits. This work was initiated by a working group between these researchers at the Santa Fe Institute, and will continue to be independently fostered by working groups among these researchers at German Centre for Integrative Biodiversity Research.

Trait evolution will follow a continuous Markov process (see Table 1) for continuous traits, and be thresholded [80] for discrete traits. Trait-environment interactions will depend on the distance between an individual’s trait value and the optimal trait value for the environment, the greater the distance the less the fitness of that individual. Trait-based species-species interactions will be again based on the distance between traits, in this case the closer the trait values of two individuals, the stronger their Lotka-Volterra [81] competition. This full model (Fig. 2, Table 1) will jointly predict species abundance data, genetic and phylogenetic data, and trait distributions.

Because available population genetic data across many individuals in a community are limited to extremely reduced genomic regions, and often mined for synonymous SNPs, we will initially model neutral mutations; however, demography will be driven by potentially non-neutral ecological dynamics, and thus population genetic measures, despite being mutationally neutral, may still be far from the predicted drift-mutation balance of a constant population.

Analytical solutions to such models might be available, but only in the asymptotic limit of equilibrium [e.g., 15, 22]. Because we are interested in potentially non-equilibrium histories (hypothesis **H5**) as a driver of biodiversity patterns, we will use efficient methods to simulate model predictions under different scenarios (from neutral to trait-based assembly) and use a combination of approximate Bayesian computation (ABC) and machine learning methods to infer which models with what specific parameter values are best supported by the data. Model fitting is discussed in section 2.6.

Figure 2 summarizes the simulation model which incorporates two key time scales and spatial scales. Across short timescales, slow rates of macroevolution and very large population sizes render diversity at the largest spatial scale, the regional pool, approximately constant. Abundances in the much smaller local community fluctuate due to a birth-death-immigration process [82], modified under the non-neutral models to include trait-based interactions (Table 1). The fluctuating local population determines the effective population size N_e which, together with mutation rate, determines the distribution of genetic diversities across species. Across long time scales new species arise and traits evolve.

The assembly of local communities by birth, death, and immigration is well studied and we will use these results to implement efficient forward time simulations [36, 82]. These forward-time assemblage histories will feed into aggregate population genetic histories that can be efficiently simulated via extensions of the Kingman coalescent [83]. We will use msprime [84] to implement this coalescent approach to simulating population genetic data. These simulations depend on the time trajectory of effective population size (N_e) for each species and the mutation rate μ . The time trajectory of N_e for a species is determined by that species’ simulated population fluctuations [85], which in turn derive from the parameters of birth, death, and immigration, which are fit to the data. Mutation rate μ is a free parameter that will also be fit to the data.

2.6 Model fitting

Because we are interested in the non-analytically tractable non-equilibrium dynamics of eco-evolutionary community assembly (**H5**), we will use simulation-based approaches to model fitting, including new advances in machine learning and hierarchical ABC. Fitting our model that jointly predicts four data axes that span the three dimensions of biodiversity under many mechanistic processes is precisely the means by which we test our hypotheses about the drivers of diversity (**R3**).

While ABC has been standard for inference of complex models [33], we are also interested in exploring machine learning approaches because they enable the efficient use of high-dimensional data and parameters without specific knowledge of the joint probability distribution. Unlike available

ABC methods that can become intractable in high-dimension applications with many summary statistics, many supervised machine learning methods perform best in such cases [86, 87]. As another strategy, we will explore machine-learning Bayesian approaches that can use the full data under a flexible Bayesian framework [88], rather than having to choose particular summary statistics.

For our ABC and machine learning approaches we will explore a range of summary statistics that contain information from all four data axes (species abundance, phylogenies, population genetics, and trait distributions) while flexibly allowing some axes to be incomplete. Incomplete data can be accounted for in two ways: (1) summary statistics pertaining to the missing data can be masked, effectively integrating over the possible values those data could take; or (2) a sampling model can be added to allow for the inclusion of the sparse data that do exist. The latter solution is particularly relevant when considering the occurrence data available from many studies. In these cases, we will simulate actual abundance data, but then coarse-grain those abundances to occurrences according to a binomial sampling process with detection probability p . Using such occupancy type models for occurrence data is not unprecedented [89]. However, in our approach we will be able to even more accurately estimate p because we will augment studies with only occurrence data with actual abundance surveys. Where occurrence points and surveys overlap in geographic regions, we will be able to precisely fit the value of p , which then will be carried by the model to all regions where occurrence data are available.

Based on preliminary work by Rominger, summary statistics based on Renyi entropies (equal to the logarithm of the Hill number) of SADs in a machine learning context can be as good or better than pure likelihood approaches in parameter estimation. We will therefore begin with Renyi entropies of SADs [90], distributions of genetic diversities across species, phylogenetic diversity [91], and trait distributions as our summary statistics. We will evaluate whether Renyi entropies are sufficient or if additional statistics improve fit.

2.7 Hypothesis testing

Fitting our model to real data from diverse ecosystems will allow us to understand whether a common set of mechanisms can be seen as universally driving patterns of diversity (**R3**). By evaluating how the best-fit model parameters vary across taxa and systems we will test our five main hypotheses. Figure 3 shows how correlations among estimated parameter values and between parameter values and intrinsic characteristics of regions and taxa can confirm or refute our hypotheses. Critically, not all systems are represented by the same taxa, nor do they cover the same spatial scale. In searching for correlations between model parameters we are prepared to account for random effects of taxon and study.

Evaluating whether our five biogeographic regions have achieved equilibrium (**H5**) requires added analytical nuance. For that, we will take two approaches. In the case of Hawaii and Palau, the age of ecosystems is knowable based on the geology of the system. We will thus explore explicitly how the processes inferred by our model themselves change through time, and how the community as a whole approaches, but potentially does not reach, equilibrium. In the case of Hawaii this is achieved by looking across the chronosequence [53]. In the case of Palau this is achieved by looking explicitly at the sediment core from which we will extract taxonomic, functional, and genetic data.

The proportion of equilibrium achieved [sensu 79], will also be a parameter in our simulation models (i.e. it will determine for how long a given simulation is run). By allowing this parameter to be fit to all our datasets, not just those from Hawaii and Palau, we can understand whether the community of interest has reached equilibrium. If we estimate that it has not, we can then calculate what equilibrium diversity would be under the best fit model. Comparing these observations to the real data will allow us to understand whether the non-equilibrium conditions produce higher or lower diversity across the four data axes that map onto the three dimensions of biodiversity, and

whether the specific assembly process is predictive of diversity overshoots or undershoots relative to equilibrium (**H5a-b**).

3 Broader Impacts

Our proposed project contributes to broader society in three quite different ways: (1) training and mentoring in data science, including specifically for underrepresented groups; (2) contributing open source software for science and conservation; and (3) adding to the Tree of Life.

3.1 Training and mentoring

In partnership with Data Carpentry we will develop new workshop content for the analysis of heterogeneous biodiversity data. Data Carpentry is a non-profit organization that teaches hands-on, domain specific, and evidence-based workshops. These workshops provide scientists the core computational and analytical skills they need to work with large and complex datasets. We will organize a lesson development hackathon to assemble the content of the workshop with the PIs, Co-PIs, and members of the Data Carpentry staff and community. We will use this content in a national workshop for ≈ 30 participants to be hosted by the co-PIs at the University of Florida, who have organized 10+ Data and Software Carpentry workshops. The data science principles taught in the workshops will be practiced by students—who will receive support to attend—as they prepare, curate, and publish the new data that will be generated for our proposed research.

To make this educational material as broadly accessible as possible, the biodiversity data science curriculum created in the course of this workshop will be developed as a massive open online course (MOOC) hosted on the Santa Fe Institute’s online education platform. This content will be licensed under creative commons. The creation of this course will leverage existing infrastructure and expertise at the Santa Fe Institute, including the SFI video recording studio, a robust online education platform and learning management system, and the expertise of the SFI online education staff and media specialists.

3.2 Open source software for science and conservation

3.3 Informatic platform

We will provide open-source data management software that fills the gaps of existing biodiversity data management and sharing tools, developing a powerful and intuitive software that can help research groups manage, curate, share, and publish their diverse data streams associated with large, multi-dimension, multi-institution projects exemplified by the DoB projects. This tool will also allow multiple user experiences, from programmatic interface through an R package, to a web-based graphical user interface enabled through Shiny apps.

3.4 Eco-evolutionary synthesis modeling

Our modeling framework, including simulations under different model specifications, model fitting, and model selection, will be made available as an open source R package. With this modeling power at hand, users will be able to gain completely new insights into their study systems. Whether or not communities are assembled primarily by in situ evolution or ex situ immigration strongly determines their response to anthropogenic pressures [92] and optimal conservation management [93]. We will highlight this importance in an open access applications manuscript which we will seek to publish in a conservation-oriented journal. We will specifically seek publicity for this work such that its message, and our open source software, can be brought to the attention of conservation practitioners.

3.5 Adding to the tree of life

Expanding the Tree of Life is a community-wide priority, as we ought to document diversity before human actions erase some of the most remarkable realizations of the evolutionary process. The majority of our data collection is aimed at adding robust tips to the Tree of Life, complementing other ongoing projects [61].

4 Timeline and Project Management

In the table below we detail the timeline for major milestones of our project.

Activity	Yr 1	Yr 2	Yr 3	Yr 4	Yr 5
Organizational working group at Santa Fe Institute (SFI)	■		■		■
Informatics pipeline	■	■	■		■
Data Carpentry content hackathon	■	■			
Data Carpentry workshop		■			
SFI MOOC development			■		
Filling data gaps (abundance, genetics, traits)	■	■	■		
Model development	■	■	■		
Model validation and fitting with simulated data	■		■	■	
Model-based inference with real data				■	■
Capstone workshop at SFI					■

4.1 Project Team Coordination

We have assembled a diverse team of researchers from theoreticians to systematists to field ecologists. Rominger will lead coordinating efforts across institutions remotely, and by bringing together the group for logistical and synthesis meetings at the Santa Fe Institute. Remote meetings will be held at least monthly to discuss progress on independent tasks and on synthesis. Our larger group will be organized into three core teams: (1) systems data gathering team, (2) informatics team, and (3) theory development team. The **systems team** will compile existing and collect new data from our five biogeographic regions: Rominger and Gillespie (Hawaii), Dawson (Palau), Carnaval and Michelangeli (Atlantic Forest), Owens and Robinson (Amazonia), and Owens and Soltis (US-China)), along with their collaborators. The **informatics team** (Rominger, Guralnick, and Owens with support for a scientific programmer to be housed at University of Florida) will coordinate with both the systems and theory teams to build an informatics platform that satisfies both groups' needs. The informatics team will also be responsible for developing the biodiversity data science training curriculum in collaboration with Data Carpentry and the Santa Fe Institute Office of Education. The **theory team** (Rominger, Hickerson, and collaborators Chase and Harmon) will build our novel modeling framework that jointly predicts species abundances, population and phylo-genetics, and trait distributions. They will also coordinate with the systems team to ensure the model's biological realism.

5 Results from Prior Funding

Carnaval. DEB 1343578: Dimensions US-BIOTA-Sao Paulo: A multidisciplinary framework for biodiversity prediction in the Brazilian Atlantic forest hotspot; 2013–2018; \$1,991,480. **Intellectual Merits.** The team is developing models of spatial patterns of biodiversity in the megadiverse Atlantic Forest of Brazil. They generate and integrate: (1) environmental data from novel remote sensing-based datasets and paleoenvironmental archives; (2) locality, phylogenetic, genomic, and trait data from plants and animals; and (3) new methods that incorporate big genomic and environmental data into predictive models. The project has generated 80+ publications to date, including high impact journals [2, 94–101]. **Broader Impacts.** The project contributed to training of five Ph.D. students (2 female, 2 Latin-American), 11 M.Sc. students (10 female, 9 underrepresented minority), 7 undergraduates (all female/URM), and 5 post-docs (2 female, 2 Latin-American).

Dawson. OCE-1243970: Dimensions: Collaborative Research: Do parallel patterns arise from parallel processes; 2013–2017 plus 1-yr no cost extension; \$1,369,982. **Intellectual Merit:** We investigated taxonomic, genetic, and functional diversity of microbes and macrobiota in widespread yet under-explored marine lake ecosystems, which present independent eco-evolutionary “natural experiments.” We are revealing how parallel processes in space and time influence genetic, phylogenetic, and functional diversity. We surveyed over 14000 points across 15 lakes in Palau, sampling 9000 specimens, and barcoding 1350 of invertebrates and algae. We are finding (1) marine lakes conform with island theory, (2) microbes and macrobiota can share similar biogeographic patterns, (3) abiotic lake dynamics are driven by regional climate and local factors, and (4) life-history and chance strongly influence population structure in invertebrates. Currently available research products: [102–107]. Published datasets: BCO-DMO Project 2238 datasets. **Broader Impacts:** 4 female graduate students (incl. 1 Pacific islander), 2 female lab techs, and a citizen scientist trained on the project. Advised local resource managers on conservation.

Gillespie : DEB 1241253 Dimensions: A community level approach to understanding speciation in Hawaiian lineages. 2013–2019; \$1,181,407. **Intellectual Merit.** This project aims to integrate evolutionary and ecological theory to transform understanding of the dynamics underlying community assembly. The synergy between the two approaches is made possible through the use of a habitat chronosequence, provided by the dynamic geomorphology of the young islands of the Hawaiian archipelago. We selected 6 replicates in each of 12 sites and are processing thousands of arthropod specimens while creating an mtDNA barcode library and testing metabarcoding approaches. From these data we are estimating macroecological metrics and conducting food web analysis. For focal lineages, genomic data is providing information on population differentiation over the island chronology. Papers to date: [53, 71, 108–118]. **Broader Impacts.** Research is integrated into education; trained 9 postdocs, 10 graduate students, 44 undergraduates, and one high school student.

Guralnick : DEB-1262610: Collaborative Research: ABI Development: Advancing Map of Life’s Impact and Capacity for Sharing, Integrating, and Using Global Spatial Biodiversity Knowledge; \$560,861; 2013–2018 plus one year no-cost extension. **Intellectual Merit:** The Map of Life project focuses data integration based from a variety of sources, including biocollections, species attribute data, expert assessment of range boundaries, surveys, and inventories. This work strongly enhances methods to publish, discover, openly link, and create new knowledge, in order to further assessment of global biodiversity. Papers to date: [62, 119–124]. **Broader Impacts.** The support has yielded web platforms that provide data and knowledge and provided support for 2 Ph.D. students.

Hickerson : DEB 1253710: Career: Dynamic models of isolation and admixture for community-scale population genomic inference; 2013–2017; \$667,074 plus 1-yr no cost extension. **Intellectual Merits.** Hickerson and colleagues developed inference methods to infer the aggregate spatial history of species assemblages given genome-wide data from multiple taxa. The grant has supported the publication of 22 peer reviewed articles [125–146]. **Broader Impacts.** The grant has supported two PhD students, two postdocs, one undergraduate researcher who has gone on to pursue a PhD in population genetics and an extensive educational outreach component.

Owens. DBI 1523732: Postdoctoral Fellowship in Biology: Out of the Tropics and Out of the Drawer: Integrative Analysis of the Tropical Diversity Gradient from Museum Collections of New World Swallowtail Butterflies; 2015–2017; \$138,000. **Intellectual Merit.** I imaged and processed over 1300 butterfly specimens housed at four institutions through a digitization pipeline I designed for this project. Results, leading to one publication to date [9], suggest abiotic niche conservatism and temperate-to-tropical dispersal are important, under-appreciated sources of diversity in the

tropics. **Broader impacts.** I engaged worldwide volunteers through the Notes From Nature platform to transcribe specimen label data. I have also participated in several public outreach events at the Florida Museum of Natural history to demonstrate how data derived from specimens can inform our understanding of macroecology.

Soltis. EF-1115210/DBI-1547229. Digitization HUB: A Collections Digitization Framework for the 21st Century/iDigBio Phase 2; \$12,661,986/\$15,486,747; 2011–2016 and 2016–2021. **Intellectual Merit.** iDigBio is the national coordinating center for digitization of biodiversity collections, the goal of which is to make data for millions of biological specimens available in electronic format. iDigBio currently serves over 108 million specimen records and over 22 million media records, providing information on taxonomy, geographic location, collector, date of collection, etc., and notes on phenology, habitat, molecular resources, and other features, as well as images and vocalizations. These diverse data promote integrative biodiversity research, contain untapped trait data, and provide an immense baseline for assessing impacts of climate change, invasive species, and other environmental issues. The project has produced over 20 papers to date, including in *Nature*, *BioScience*, *Proceedings of the IEEE International Conference on e-Science*, *New Phytologist*, *Philosophical Transactions of the Royal Society B*, *American Journal of Botany*, *Cladistics*, and *ZooKeys*, among others **Broader Impacts.** More than 100 workshops and hackathons have fostered training in digitization workflows, use of digitized data in research, applications in citizen science, education, and outreach. Soltis has trained 2 post-docs, 6 graduate students, and 5 undergraduates.